


# Writing Performant Concurrent Data Structures

Adrian Alic


Software Engineer @ DFINITY  
Website: <https://alic.dev>  
Contact: [contact@alic.dev](mailto:contact@alic.dev)

Rust Meetup Zürich  
March 28, 2023




# Overview

Case-study: Multi-producer, single-consumer queue.



# Overview

Case-study: Multi-producer, single-consumer queue.




Goals:

- How to write such a queue
- How to make it fast
- How to reason about correctness


# MOTIVATION

# A Multi-Core Logger



**Figure:** A sketch of a 5-core RISC-V SoC.


# The Problem With Locks



**Figure:** Locking causes unpredictable latency jitter.

# THE IDEA

# A Bunch of Ring Buffers



concurrent threads produce  
data stream

push data to thread-local  
FIFO ring buffer

consumer polls and empties  
the queues in a loop



# Naive Rust Definition

```
// if you like pointer indirection
```

```
struct TLQ {  
    buffer: Vec<u8>,  
    head: u16,  
    tail: u16,  
}
```

```
// if buffer size is known at compile-time
```

```
struct TLQ<const C: usize> {  
    buffer: [u8; C],  
    head: u16,  
    tail: u16,  
}
```

**However:** this definition has some problems...


# Lack of Cache Locality

If we store *multiple* TLQs in an array, iterating over heads and tails becomes costly.



# Lack of Cache Locality

If we store *multiple* TLQs in an array, iterating over heads and tails becomes costly.



This problem of traversing fields is common in game development (ECS).

# Improving Cache Locality

One solution: *Struct of Arrays*.

# Improving Cache Locality

One solution: *Struct of Arrays*.

```
struct Offset {  
    head: u16,  
    tail: u16,  
}
```

# Improving Cache Locality

One solution: *Struct of Arrays*.

```
struct Offset {  
    head: u16,  
    tail: u16,  
}
```

```
struct Buffer<const C: usize> {  
    buffer: [u8; C],  
}
```

# Improving Cache Locality


One solution: *Struct of Arrays*.

```
struct Offset {  
    head: u16,  
    tail: u16,  
}
```

```
struct Buffer<const C: usize> {  
    buffer: [u8; C],  
}
```

```
struct Queue<const T: usize, const C: usize> {  
    offsets: [Offset; T],  
    buffers: [Buffer<C>; T]  
}
```

# New Layout Visualized



**Figure:** Our consumer can now iterate through all offsets without tons of cache misses.

Some languages like Zig have built-in support for the SoA pattern<sup>1</sup>.

<sup>1</sup><https://kristoff.it/blog/zig-multi-sequence-for-loops/>



# THE MEMORY MODEL

# The Illusion of Safety on x86



**Figure:** Don't do this. The memory ordering I chose for my atomic ops only worked on x86, but **blew up** on a *weaker* memory model (aarch64).

# Segfaults on aarch64

Property		Alpha	Arm7-A/R	Arm8	Itanium	MIPS	POWER	SPARC TSO	x86	z Systems
Memory Ordering	Loads Reordered After Loads or Stores?	Y	Y	Y	Y	Y	Y			
	Stores Reordered After Stores?	Y	Y	Y	Y	Y	Y			
	Stores Reordered After Loads?	Y	Y	Y	Y	Y	Y	Y	Y	Y
	Atomic Instructions Reordered With Loads or Stores?	Y	Y	Y		Y	Y			
	Dependent Loads Reordered?	Y								
	Dependent Stores Reordered?									
	Non-Sequentially Consistent?	Y	Y	Y	Y	Y	Y	Y	Y	Y
	Non-Multicopy Atomic?	Y	Y	Y	Y	Y	Y	Y	Y	
	Non-Other-Multicopy Atomic?	Y	Y		Y	Y	Y			
	Non-Cache Coherent?				Y					

**Figure:** McKenney [1, p. 352] lists differences between hardware platforms in detail.

# C11 Memory Model

Rust follows the C11 memory ordering spec<sup>2</sup>. It includes:

---

<sup>2</sup>[https://en.cppreference.com/w/cpp/atomic/memory\\_order](https://en.cppreference.com/w/cpp/atomic/memory_order)

# C11 Memory Model

Rust follows the C11 memory ordering spec<sup>2</sup>. It includes:

Specification of *modification order*:

- RR/RW/WR/WW Coherency

Flavors of "before":

- Sequenced-before
- Dependency-ordered before
- Inter-thread happens-before
- Happens-before


Also relevant: *evaluation order*<sup>3</sup>

---

<sup>2</sup>[https://en.cppreference.com/w/cpp/atomic/memory\\_order](https://en.cppreference.com/w/cpp/atomic/memory_order)


<sup>3</sup>[https://en.cppreference.com/w/cpp/language/eval\\_order](https://en.cppreference.com/w/cpp/language/eval_order)

# Concurrency Behavior of Our Queue



<sup>4</sup>[https://doc.rust-lang.org/std/sync/atomic/struct.AtomicU64.html#method.compare\\_exchange](https://doc.rust-lang.org/std/sync/atomic/struct.AtomicU64.html#method.compare_exchange)

# Concurrency Behavior of Our Queue



Our queue is essentially an SPSC without competing stores - thus we have no need for atomic RCU primitives<sup>4</sup>.

<sup>4</sup>[https://doc.rust-lang.org/std/sync/atomic/struct.AtomicU64.html#method.compare\\_exchange](https://doc.rust-lang.org/std/sync/atomic/struct.AtomicU64.html#method.compare_exchange)

# The Two Basic Queue Operations

Our SPSC requires two release-acquire pairs. We can look at the first one below.

```
// producer thread
fn push(data) {
    h = head.load(_)
    new_h = h + data.len()

    // write data
    buffer[h..new_h] = data;

    // update index
    h.store(new_h, _)
}
```

```
// consumer thread
fn pop() [u8] {
    // read index
    h = tail.load(_)
    t = tail.load(_)

    // read data
    buffer[t..h]
}
```



# The Two Basic Queue Operations

Our SPSC requires two release-acquire pairs. We can look at the first one below.

```
// producer thread
fn push(data) {
    h = head.load(_)
    new_h = h + data.len()

    // write data
    buffer[h..new_h] = data;

    // update index
    h.store(new_h, release)
}
```


```
// consumer thread
fn pop() [u8] {
    // read index
    h = tail.load(acquire)
    t = tail.load(_)

    // read data
    buffer[t..h]
}
```

# IMPLEMENTATION IN RUST

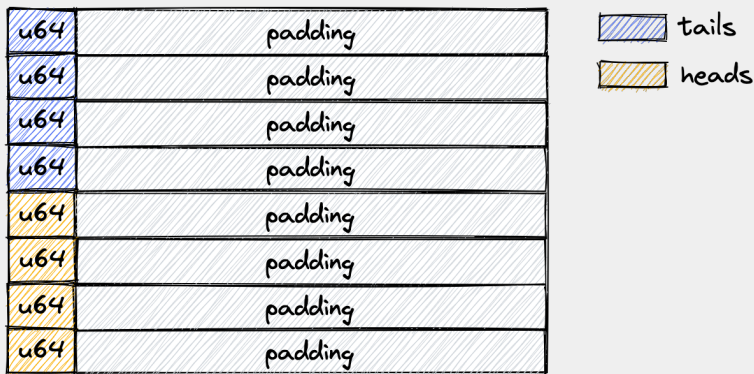
# Avoiding False Sharing

Since offsets are accessed concurrently, we need to be aware of cache coherence effects.




**Figure:** The most common solution is to pad all shared fields to a cache line.

# Cache-Alignment for Each Offset



**Figure:** Fully padded version. No false sharing will occur.

# A Possible Middle Ground



**Figure:** This hybrid version allows for atomic batch updates.

# Implementation in Rust

```
#[repr(align(64))]  
struct Tail(u16);
```


```
#[repr(align(64))]  
struct Head(u16);
```

```
struct Offsets<const T: usize> {  
    tails: [Tail; T],  
    heads: [Head; T],  
}
```

// Or alternatively, use the crossbeam\_util crate

```
struct Offsets<const T: usize> {  
    tails: [CachePadded<Tail>; T],  
    heads: [CachePadded<Head>; T],  
}
```


# False Sharing Can Have a Large Impact



**Figure:** From a benchmark on false sharing <sup>5</sup>

<sup>5</sup><https://alic.dev/blog/false-sharing>


# Consumer-Side Pointer Compression



**Figure:** We can decrease the addressing granularity, reducing memory footprint.



# Pointer Compression Visualized



# Implementation in Rust

```
struct Consumer<const C: usize> {
    shared_tail: *const AtomicU16,
    local_tail: usize,
}


fn update_tail(&mut self, val) {
    self.local_tail = val;
    self.shared_tail.store(
        compress(self.local_tail, C), // <---
        Ordering::Release
    );
}

fn compress(tail: usize, C: usize) -> u16 {
    let shift = if C <= 16 { 0 } else { C - 16 };
    (tail >> shift)
}
```

# Local Caching of Offsets

# Local Caching of Offsets

 tail     head     elements



# Local Caching of Offsets

 tail     head     elements



# Local Caching of Offsets

 tail     head     elements



# Local Caching of Offsets

 tail     head     elements



# **CRAFTING SAFE ABSTRACTIONS**



# Limits of the Borrow Checker

The borrow checker and lifetime system is not designed to reason about correctness of arbitrary concurrent data structures.

Example: *Atomics*

```
impl AtomicUsize {  
    pub fn store(&self, val: bool, order: Ordering) {  
        // SAFETY: any data races are prevented by atomic  
        // intrinsics and the raw pointer passed in is  
        // valid because we got it from a reference.  
        unsafe {  
            atomic_store(self.v.get(), val as u8, order);  
        }  
    }  
}
```

# Newtyping Heads and Tails

Newtyping your data structures to give them semantics can prevent many subtle bugs.

```
type utail = u16;  
type udefault = u32;
```

```
type AtomicTail = AtomicU16;  
type AtomicHead = AtomicU32;
```

```
// Read and write permissions
```

```
struct RWHead<const C: usize>(*const AtomicHead);  
struct RWTail<const C: usize>(*const AtomicTail);
```

```
// Read-only permission
```

```
struct ReadOnlyHead<const C: usize>(*const AtomicHead);  
struct ReadOnlyTail<const C: usize>(*const AtomicTail);
```

# Incorporating Newtypes Into Data Structure

Good newtypes communicate intent *clearly*.

```
pub struct Consumer<...> {  
    tails: [RWTail<C>; T],  
    heads: [ReadOnlyHead<C>; T],  
    buffer: ReadOnlyBuffer<T, S, L>,  
}
```

```
pub struct Producer<...> {  
    pub head: RWHead<C>,  
    pub tail: ReadOnlyTail<C>,  
    pub buffer: RWBuffer<L>,  
}
```

## Const Generics Help With Safety

```
impl<
    const T: usize,    // # of producers
    const C: usize,    // bitsize of queue
    const S: usize,    // # of bytes (total)
    const L: usize,    // # of bytes (per producer)
    A: ThreadSafeAlloc, // custom allocator type
> ProducerHandle<T, C, S, L, A> {
    // ...
}
```

## Reading From Queue With RAI

```
fn pop(&self, pid: usize) -> Vec<u8>;
```

## Reading From Queue With RAI

```
fn pop(&self, pid: usize) -> Vec<u8>;
```

```
fn pop(&self, pid: usize, dst: &mut [u8]) -> usize;
```

## Reading From Queue With RAI

```
fn pop(&self, pid: usize) -> Vec<u8>;
```

```
fn pop(&self, pid: usize, dst: &mut [u8]) -> usize;
```

```
fn pop<'a>(&'a mut self, pid: usize) -> &'a [u8];
```

# Reading From Queue With RAI

```
fn pop(&self, pid: usize) -> Vec<u8>;

fn pop(&self, pid: usize, dst: &mut [u8]) -> usize;

fn pop<'a>(&'a mut self, pid: usize) -> &'a [u8];

fn pop<'a>(&'a mut self, pid: usize) -> Section<'a>;

struct Section<'a>{buffer: &'a [u8], ... };

impl<'a> Drop for Section<'a> {
    fn drop(&mut self) {
        unsafe {
            // increment tail atomically
        }
    }
}
```



# Reading From Queue With RAI

```
// max capacity is 2^3 - 1
let (tx, mut rx) = wfmpsc::queue!(bitsize: 3, producers: 1);
tx[0].push(b"5678901");
{
    let mut section = rx.pop(0);
    for c in section.get_buffer().iter() {
        // iterate over section and do things
    }
} // dropping buffer
```

# Reading From Queue With RAI

```
// max capacity is 2^3 - 1
let (tx, mut rx) = wfmpsc::queue!(bitsize: 3, producers: 1);
tx[0].push(b"5678901");
{
    let mut section = rx.pop(0);
    for c in section.get_buffer().iter() {
        // iterate over section and do things
    }
    let mut another_one = rx.pop(0);
    //          ^^^^^^^^^
    //          |
    //          + can't create another section
    //          while previous one in scope
    black_box(&section);
} // dropping buffer
```

# **RUNTIME ANALYSIS WITH MIRI**

# What Is Miri?

**Miri**<sup>6</sup> is an interpreter for Rust's Mid-Level IR that dynamically checks for undefined behavior.

Checks include:

- OOB memory access & use-after-free
- Illegal memory alignments
- Reading from uninitialized memory
- Data races
- Violation of stacked borrows aliasing model

---

<sup>6</sup><https://github.com/rust-lang/miri>

# Issue #1: Uninitialized Arrays

Can you spot a potential problem here?

```
let mut producers: [Producer<...>; T] = { mem::zeroed() };

for (i, p) in producers.iter_mut().enumerate() {
    *p = self.get_producer_handle(i);
    //      ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
    //      |
    //      assume this function returns a valid object
}
```

# Issue #1: Uninitialized Arrays

Can you spot a potential problem here?

```
let mut producers: [Producer<...>; T] = { mem::zeroed() };

for (i, p) in producers.iter_mut().enumerate() {
    *p = self.get_producer_handle(i);
    //      ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
    //      |
    //      assume this function returns a valid object
}
```


**Problem:** The assignment calls `Drop::drop` on the old value. This violates the producer's atomic refcount invariant.

# Issue #1: Uninitialized Arrays

```
let mut producers: [MaybeUninit<Producer<...>>; T] =
    unsafe { MaybeUninit::uninit().assume_init() };

for (i, p) in producers.iter_mut().enumerate() {
    p.write(prod_handle(ptr, i as u8));
}
// FIXME: Cannot do mem::transmute from MaybeUninit to
// a const generic array.
// See https://github.com/rust-lang/rust/issues/61956
let prod_ptr = addr_of!(producers) as *const _;
let producers = unsafe { core::ptr::read(prod_ptr) };
```

## Issue #2: Dangling Pointer



**Figure:** Elements can spill over the boundary of the ring buffer, so we need to invoke memcpy twice.



## Issue #2: Dangling Pointer

```
// first memcpy
core::ptr::copy_nonoverlapping(
    src as *const u8,
    dst as *mut u8,
    L - head,
);
// second memcpy
core::ptr::copy_nonoverlapping(
    (src + C - head) as *const u8,
    self.buffer.0 as *mut u8,
    len - L + head,
);
```

## Issue #2: Dangling Pointer

```
// first memcpy
core::ptr::copy_nonoverlapping(
    src as *const u8,
    dst as *mut u8,
    L - head,
);
// second memcpy
core::ptr::copy_nonoverlapping(
    (src + C - head) as *const u8,
    self.buffer.0 as *mut u8,
    len - L + head,
);
```

## Issue #3: Incorrect Pointer Arithmetics (again)

error: unsupported operation: racy imperfectly overlapping  
atomic access is not possible in the C++20 memory model,  
and not supported by Miri's weak memory emulation

--> /Users/zk/wfmpsc/src/lib.rs:275:13

```
275 |         atomic.store(val, ord);  
    |         ^^^^^^^^^^^^^^^^^^^^^^^^^ racy imperfectly  
    |         overlapping atomic access is not possible  
    |         in the C++20 memory model, and not  
    |         supported by Miri's weak memory emulation
```

## Issue #3: Incorrect Pointer Arithmetics (again)



# CONCLUSION

# Conclusion

- Be cognisant of the language's semantic model

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>

# Conclusion

- Be cognisant of the language's semantic model
  - ▶ The Rustonomicon<sup>7</sup> is a good starting point

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>

# Conclusion

- Be cognisant of the language's semantic model
  - ▶ The Rustonomicon<sup>7</sup> is a good starting point
- Familiarize yourself with the memory models that underpin your stack

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>



# Conclusion

- Be cognisant of the language's semantic model
  - ▶ The Rustonomicon<sup>7</sup> is a good starting point
- Familiarize yourself with the memory models that underpin your stack
- Use RAI and lifetimes to create safe viewtypes

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>

# Conclusion

- Be cognisant of the language's semantic model
  - ▶ The Rustonomicon<sup>7</sup> is a good starting point
- Familiarize yourself with the memory models that underpin your stack
- Use RAI and lifetimes to create safe viewtypes
- Memory fragmentation is a powerful trade off

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>

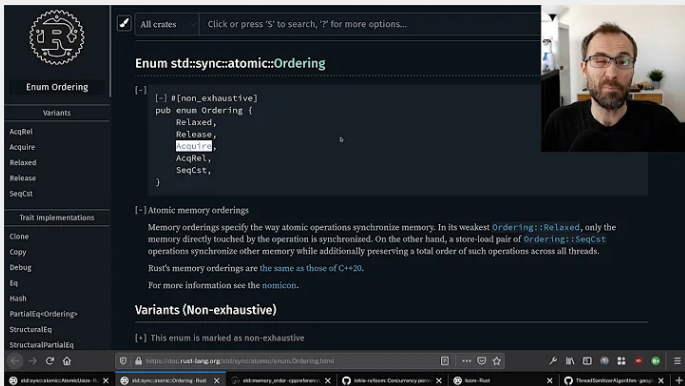
# Conclusion

- Be cognisant of the language's semantic model
  - ▶ The Rustonomicon<sup>7</sup> is a good starting point
- Familiarize yourself with the memory models that underpin your stack
- Use RAI and lifetimes to create safe viewtypes
- Memory fragmentation is a powerful trade off
- Learn from the OGs

---

<sup>7</sup><https://doc.rust-lang.org/nomicon/>

# More Resources



The screenshot shows the Rust documentation for `std::sync::atomic::Ordering`. On the left is a sidebar with a search icon and a list of navigation links: Enum Ordering, Variants, AcqRel, Acquire, Relaxed, Release, SeqCst, Trait Implementations, Clone, Copy, Debug, Eq, Hash, PartialEq<Ordering>, StructuralEq, and StructuralPartialEq. The main content area has a search bar at the top. Below it, the title `Enum std::sync::atomic::Ordering` is followed by a code block showing the enum definition: 


```
[-] #[non_exhaustive]
pub enum Ordering {
    Relaxed,
    Release,
    Acquire,
    AcqRel,
    SeqCst,
}
```

 Below the code, there is a section for Atomic memory orderings, explaining that memory orderings specify how atomic operations synchronize memory, with `Ordering::Relaxed` being the weakest and `Ordering::SeqCst` being the strongest. It also notes that Rust's memory orderings are the same as those of C++20. A link to the `nomicon` is provided for more information. The next section is Variants (Non-exhaustive), which states that this enum is marked as non-exhaustive. At the bottom of the page, a video inset shows Jon Gjengset speaking. The browser's address bar shows the URL `https://doc.rust-lang.org/std/sync/atomic/enum.Ordering.html`. The taskbar at the bottom shows several open windows, including the Rust documentation, a terminal, and a video player.

**Figure:** Atomics and Memory Ordering by Jon Gjengset [\[video\]](#)

**THANKS FOR YOUR ATTENTION!**

# References

-  PAUL E MCKENNEY.  
**IS PARALLEL PROGRAMMING HARD, AND, IF SO, WHAT CAN YOU DO ABOUT IT?**  
*arXiv preprint arXiv:1701.00854*, 2017.