

Unsafe Rust and Miri

Ralf Jung (ETH Zurich)

Rust Zurich Meetup, June 2023

Rust is a memory-safe language.

Rust is a memory-safe language.

The compiler guarantees that Rust programs do not

- Access memory out-of-bounds
- Access memory after it has been freed
- Execute the “wrong” code
- Have data races

Rust is a memory-safe language.

The compiler guarantees that Rust programs do not

- Access memory out-of-bounds
- Access memory after it has been freed
- Execute the “wrong” code
- Have data races

Except...

unsafe code

```
fn evil() { unsafe {  
    let x = 0i32;  
    let ptr = ptr::from_ref(&x);  
    ptr.offset(1).read(); // uh-oh...  
} }
```

In unsafe Rust, memory safety
becomes the programmer's
responsibility.

The Contract

There is a contract between programmer and compiler, with obligation for both sides

- **Compiler's obligations:**
Generate a binary that matches the Rust source code
- **Programmer's obligations:**
Write a program that follows “the rules”

If you break that contract, **anything can happen!**



If you break that contract, **anything can happen!**



We call this **Undefined Behavior.**

But what about “soundness”?

Soundness is the property of an API to **avoid Undefined Behavior when called from safe code.***

```
pub fn call_me(x: &bool) { unsafe {  
    let ptr = ptr::from_ref(x).cast::<u8>();  
    if ptr.read() == 2 {  
        hint::unreachable_unchecked();  
    }  
} }
```

But what about “soundness”?

Soundness is the property of an API to **avoid Undefined Behavior when called from safe code.***

```
pub fn call_me(x: &bool) { unsafe {  
    let ptr = ptr::from_ref(x).cast::<u8>();  
    if ptr.read() == 2 {  
        hint::unreachable_unchecked();  
    }  
} }
```

What about the **raw pointer version**?

Key terminology

- **Undefined Behavior:** A program that violates “the rules” during its execution
- **Unsound Code:** A library that can be used by safe code to cause Undefined Behavior

Common misconceptions

- “My code is fine, I tested it.”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”
- “This code was fine with a previous compiler, but now it broke; this is the compiler’s fault.”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”
- “This code was fine with a previous compiler, but now it broke; this is the compiler’s fault.”
- “That evil compiler should tell me about the UB instead of exploiting it!”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”
- “This code was fine with a previous compiler, but now it broke; this is the compiler’s fault.”
- “That evil compiler should tell me about the UB instead of exploiting it!”
- “This is terrible, I will go back to C / use inline assembly.”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”
- “This code was fine with a previous compiler, but now it broke; this is the compiler’s fault.”
- “That evil compiler should tell me about the UB instead of exploiting it!”
- “This is terrible, I will go back to C / use inline assembly.”
- “I disabled optimizations so UB cannot be a problem.”

Common misconceptions

- “My code is fine, I tested it.”
- “It’s what the hardware does”
- “This code was fine with a previous compiler, but now it broke; this is the compiler’s fault.”
- “That evil compiler should tell me about the UB instead of exploiting it!”
- “This is terrible, I will go back to C / use inline assembly.”
- “I disabled optimizations so UB cannot be a problem.”
- “It’s fine if I put `unsafe` around it.”

Knowing “the rules” is
absolutely crucial
when writing
unsafe code!

But knowing all the rules is hard...

Thankfully, there is a tool to help:

Miri

But knowing all the rules is hard...

Thankfully, there is a tool to help:

Miri

- Miri can detect Undefined Behavior in a program or test suite
- Miri cannot ensure a library is sound!

But knowing all the rules is hard...

Thankfully, there is a tool to help:

Miri

- Miri can detect Undefined Behavior in a program or test suite
- Miri cannot ensure a library is sound!

Caveat: Miri only supports pure Rust code.
No C FFI.

But knowing all the rules is hard...

Thankfully, there is a tool to help:

Try it on your crate:

```
rustup toolchain install nightly --component miri  
cargo +nightly miri test
```

program or test suite

- Miri **cannot** ensure a library is sound!

Caveat: Miri only supports pure Rust code.

No C FFI.

Miri cannot ensure a library is sound!

So, it is still worth knowing “the rules”.

Examples of Undefined Behavior

- Out-of-bounds access
- Use-after-free
- Out-of-bounds pointer arithmetic
- Insufficient alignment (1) (2)
- Invalid value (1) (2) (3)
- Violation of reference aliasing rules (1) (2) (3)
- Data race

Examples of Undefined Behavior

- Out-of-bounds access

- Use after free

For a more complete list, see [the reference](#). However, specifying this in full precision is still work-in-progress.

`unsafe` library functions always [document](#) their requirements, make sure to take good notice of that.

- Violation of reference aliasing rules (1) (2) (3)

- Data race

When writing unsafe Rust:



Remember not to break “the rules”!

Further reading

- With Undefined Behavior, Anything is Possible
- Undefined Behavior deserves a better reputation
- Why even unused data needs to be valid
- "What The Hardware Does" is not What Your Program Does: Uninitialized Memory